

Companion Enterprise

One AI assistant for the whole company, on our terms

■ 6 October 2026 ■ Working title ■ Builds on the Local LLM PC Companion

01 The goal

One AI assistant that everyone in the company can use for everyday work: asking questions, getting answers from our own documents with the sources named (RAG), writing and reviewing code, and working with Jira and Outlook. It runs on AI models we host ourselves, or on approved outside providers such as OpenAI, Anthropic and Google, always under the company's rules and with a record of how it is used.

02 Where it comes from: the Local LLM PC Companion

The Local LLM PC Companion is a working proof of concept built in-house. It runs open-weight AI models entirely on a laptop. It offers chat, answers questions about code, and has an agent that reads and writes project files, runs commands and checks its own work. Permission modes decide what the agent may do without asking, memory is managed automatically on long tasks, and everything is logged.

91

tokens per second from a mid-size model on a laptop graphics card: many times faster than anyone reads

3 s

to read a 12-page document

13

tokens per second with no graphics card at all

0

bytes sent outside the machine while web search is off

Measured on a laptop with a 12 GB graphics card, September 2026.

It proves that the core works, on modest hardware. What it lacks is what a company needs around it: today it serves one person, runs one model at a time, and lives on one PC.

03 What the enterprise version adds

Company sign-in. People use their Microsoft 365 account; teams and roles come from the company directory.

Many people, many models. Several models served to everyone at once; servers are added or removed as demand changes.

Our models or outside providers. Open-weight models on our own hardware, or OpenAI, Anthropic and Google through one gateway, behind the same rules.

Answers from our documents (RAG). The assistant first finds the relevant passages in company, team or personal document collections, then answers from them and names its sources; each person only gets passages from documents they are allowed to open.


Jira and Outlook. Each person connects their own account; sending an email or changing a ticket always asks first.

Tools that IT controls (MCP). Plug-in tools are MCP servers, the open standard for connecting AI to other systems; IT adds or removes them and gives them to the teams that need them, and the system checks that a model can actually use a tool before offering it.

Guards and an audit log. Rules on what may leave the company, approval before actions, spending limits, and a tamper-evident record of who asked what.

A desktop app and an admin dashboard. An installable app for everyone, and a separate dashboard for the team that runs the service.

04 How it fits together



People use the desktop app (or a browser). Company servers handle sign-in, rules, the AI models and the record of use, and grow by adding copies. Our own models run in pools on our GPUs; outside providers and services are reached only through the same rules. The green dashed line is the link that lets the coding agent work on files on the user's own PC, inside folders the user allowed.

05 From proof of concept to company service

	Local LLM PC Companion (today)	Companion Enterprise
People	One person	Everyone, with company sign-in, teams and roles
Models	One at a time, on that PC	Several at once, shared, scaling with demand
Where it runs	One laptop	Our own servers or AWS; a single server for the pilot
AI providers	Our own models only	Our own models, plus OpenAI, Anthropic and Google through one gateway
Documents	Keyword search in the files of one project	RAG over company, team and personal collections, with access rights and citations
Tools	Built in	Plug-in tools (MCP servers) managed by IT per team, including Jira and Outlook
Records	Logs on the PC	A tamper-evident audit log with retention rules
App	A page in the browser	A desktop app, plus an admin dashboard

06 The plan

Foundations

Company sign-in, users and teams, a shared database, several models at once, outside providers.

Everyday use

The desktop app, plug-in tools (MCP), Jira and Outlook.

Pilot

RAG: answers from that department's documents. One department, 20 to 30 people, on one server.

Company-ready

Guards and audit complete, many servers with automatic scaling, a security review.

Each step ends with a review before the next one starts. **Where we are (6 October 2026):** the Companion has been copied as the base with all its existing checks passing, and every part of its code has been mapped onto this design. Time estimates follow from that map.

07 What we need from the company

- A Microsoft Entra ID app registration, for sign-in and Outlook
- Accounts with the outside providers we want to offer
- One GPU server (or the AWS equivalent) and a pilot department
- A code-signing certificate for the desktop app
- A decision on hosting: our own servers or AWS